

Attention Is All You Need: Revisiting Transformer Efficiency

Abstract

We propose an efficient Transformer variant achieving 92.3% accuracy on GLUE benchmark with 30% fewer parameters.

Introduction

Self-attention models dominate NLP but suffer from quadratic complexity.

Method

Our approach uses sparse attention patterns and linear projections across layers.

Experiment

Results: GLUE 92.3%, SQuAD F1 88.7%, inference speedup 2.1x over baseline.

Conclusion

Efficient attention enables scalable deployment on edge devices.